

Encounter Based Multi Robot Simultaneous Localization and Occupancy Grid Mapping

RB Choroszuca¹, C Hyman², and A Collier³

Abstract—In this paper we present and implement the algorithm in Howard’s 2006 paper: “Multi-robot simultaneous localization and mapping using particle filters.” We also make a small modification to allow for each robot to have its own particle filter and qualitatively demonstrate that it performs better than the current state of the art.

I. INTRODUCTION

The automated exploration of unknown environments has become one of the foremost challenges in mobile robotics. For a robot to explore an environment, it must map the environment and concurrently localize itself within the environment. The framework used to perform this task is known as simultaneous localization and mapping (SLAM) and has been well covered in the literature using a variety of techniques [?], [?].

While SLAM is well known and has a rich history of successes using a single robot, it can often be a slow process due to both constraints on the robot, such as speed and data processing, and lack of redundancy, i.e. robot failure [?], [?]. To address the speed of mapping and to add redundancy, coordinated or multi-robot SLAM (MRSLAM) was created.

MRSLAM is as it sounds, SLAM using multiple exploring robots. This approach allows for the partition of the physical search space using different robots, typically decreasing the time it takes to map an area, and increasing the likelihood of full map coverage in the event of robot failure. This temporal exploration parallelism does, however, come at the cost of added complexity. The added complexity of MRSLAM comes in two major components: coordination of exploration using multiple search agents and merging the maps of these agents [?]. Coordinated exploration consists of planning the groups’ search path, most often formulated to cover the most search space in minimum time or to optimize some other mapping cost criteria. The other major component, map merging, is the combination of individual robot’s observations and maps into one cohesive global map estimate.

In this paper we use the map merging algorithm presented in Howard’s 2006 paper: “Multi-Robot Simultaneous Localization and Mapping using Particle Filters” [?]. This paper is of particular interest as it presents a resource efficient online solution to the MRSLAM map merging problem given unknown initial robot poses.

The algorithm is an encounter-based algorithm. Given that we start with one robot mapping (call this robot 1), when it encounters another robot (robot 2), the relative pose between the two robots is computed and robot 2’s odometry and scan data is transmitted in reverse order to robot 1 and integrated to create a single map posterior.

Fig. ?? depicts the process with robot 1’s triple $(x_t^1, u_{t-1}^1, z_t^1)$, and robot 2’s triple $(x_t^2, u_{t-1}^2, z_t^2)$. When $t = 2$, robot 1 observes robot 2 and determines the relative pose Δ , then the past odometry and scan data is integrated in reverse order.




Fig. 1. Depiction of data integration. Taken from [?].

The paper is structured as follows: In §?? some background about MRSLAM is provided, the contribution of [?] is discussed and a modified scheme is presented. In §?? the algorithm of [?] is discussed, and in §?? experimental results are provided for a simulated example and the Albert B data set [?].

II. BACKGROUND

A. State of the Art

In the case where all relative robot poses are known, merging maps is a trivial problem using small modifications to existing SLAM techniques [?]. In the general MRSLAM problem, however, robots may start with unknown absolute and relative poses, and therefore merging of maps requires the discovery of relative relationships between different robot

¹ RB Choroszuca is with the School of Naval Architecture and Marine Engineering, University of Michigan. E-mail: riboch@umich.edu

² C Hyman is with the School of Electrical Engineering and Computer Science, University of Michigan. E-mail: hyman@umich.edu

³ A Collier is with the School of Electrical Engineering and Computer Science, University of Michigan. E-mail: amcollie@umich.edu

trajectories to build a single map. This is often a costly process, which in general can be solved for robots sharing a search space by estimating each robot's relative pose given a partial map, but this leads to exponential complexity with respect to the number of exploring robots[?]. Nevertheless, several practical algorithms exist to circumvent this naive and inefficient approach, including coarse topological matching and stitching techniques borrowed from computer vision [?].

B. Extension to Original Work[?]

Independent weights per pose: Richard

III. THE ALGORITHM

A. Overview

Howards multi-robot SLAM algorithm generates a single map and pose posterior much in the same way an occupancy grid map SLAM algorithm based on a Rao-Blackwellized particle filter (RBPF) would, but with additions to accommodate multiple robots. Mapping begins with a set consisting of a single robot with known pose and an occupancy grid map of the environment is built as this robot traverses and measures it. The ultimate goal of the algorithm is to simultaneously compute for time t , the full SLAM posterior containing all robot pose trajectories $x_{1:t}^i$ and the global map $m_{1:t}$, given only a single known initial pose and sets of measurements

$$p(x_{1:t}^1, x_{1:t}^2, \dots, x_{1:t}^M, m_{1:t} | x_0^1, z_{1:t}^1, u_{0:t}^1, \Delta_s^2, z_{1:t}^2, u_{0:t}^2, \Delta_s^3, \dots, \Delta_s^M, z_{1:t}^M, u_{0:t}^M)$$

The odometry and measurement data of all other robots is stored as it is collected, as it cannot contribute to the global map without a known relative pose linking it to the frame of the first robot. As the first robot encounters additional robots via a mutual pose observation, the newly observed robot is added to the set of mapping robots and its actions and the map is sequentially conditioned by its actions and measurements. From the point of observation, all future actions and measurements of the new robot are used to condition the map posterior, and likewise all of its previously stored actions and measurements are played back in reverse order from that point as a virtual robot. This information is then passed to a pair of new particle filters for the pose of these robots that contribute to the same global map.

As further encounters occur between any robot, real or virtual, in the mapping set and a previously unseen robot, the previously unseen robot is added to the set of mapping robots and its measured relative pose is again used to establish a reference point with which its actions and measurements can condition the map posterior. Further mutual observations of two robots already in the mapping set are ignored for the sake of simplicity. This encounter-add process is continued recursively for all robots until all stored and future odometry and measurements are exhausted or all possible mapping robots have made an encounter with the mapping set. Beyond this point, the algorithm behaves as a normal RBPF with a stacked state containing pose posteriors of all robots in addition to the map.

This algorithm relies on a number of assumptions that are requisite for it to effectively solve the MRSLAM problem. Firstly, for all explored regions to count towards the map, each mapping robot must have encountered a robot that is an element of the mapping set in order to make a fully connected graph of robot poses, and therefore maximally complete map. Additionally, each robot pose trajectory is independent of all other trajectories such that motion or observation from one robot does not affect another outside of encounter events [?].

B. Algorithm Details

Howards MRSLAM algorithm follows the basic structure of the below steps that perform data queueing, robot encounter management, and SLAM filter updates.

- 1) Queue measurement, odometry, and mutual observation (encounter) data
- 2) For all robots in the mapping set, take first element from queue and update SLAM filter
- 3) For all new encounters, split encountered queue at present time into causal and acausal queues and instantiate a new set of particles for the causal and acausal trajectory of the observed robot.
- 4) For all available acausal queues, remove last element and use it to update the SLAM filter for the corresponding acausal robot
- 5) If data includes a prior encounter with a robot not in the mapping set, instantiate a new set of particle for this robot and its acausal virtual self at its location and split the queue for that robot into a causal and acausal queues
- 6) Repeat until all data is consumed

1) *Data Queueing*: As each robot R collects data, its current odometry u_t^R and measurements z_t^R are added to a running queue data structure before processing. An encounter flag E_t^R and relative pose Δ_t^R , if available, are also added to this queue.

$$queue_{0:t}^R = append(queue_{0:t-1}^R, \{u_t^R, z_t^R, E_t^R, \Delta_t^R\})$$

The storing of this data allows for either immediate use by particle filters for robots in the mapping set, or later joining or SLAM filtering after a future encounter with a mapping robot.

2) *Encounters*: An encounter in this sense is defined as the observation of one robot by another mapping robot that produces a relative pose estimate between the robots of Δ . For this implementation, it was assumed that the estimation of Δ was deterministic for the sake of demonstration and simplicity.

Upon an encounter at time s between robot A in the mapping set with a pose x_s^A and a robot B not in the set with a measured relative pose Δ_s^{AB} , a reference frame pose for each particle i of the newly observed robot is instantiated as causal particle posterior $x_s^{B(i)}$ and acausal particle pose

posterior $\bar{x}_s^{B(i)}$ as:

$$\begin{aligned} x_s^{B(i)} &= x_s^{A(i)} \oplus \Delta_s^{AB} \\ \bar{x}_s^{B(i)} &= x_s^{A(i)} \oplus \Delta_s^{AB} \end{aligned}$$

Where \oplus is the pose composition operator. SSC reference? At this time, the stored data queue for robot B is also split into a causal queue $queue^B$ and acausal queue $queue^B$.

$$\begin{aligned} queue_{t+1}^B &= reverse(queue_{0:t-1}^B) \\ queue^B &= queue_{t+1}^B. \end{aligned}$$

If neither robot is within the mapping set and an encounter occurs, this encounter and the associated measured relative pose is stored for later use. In the event that one of these robots later encounters a robot in the mapping set, its virtual acausal robot will spawn another robot pair of the unencountered robot upon backtracking to the location of this initial encounter. At this acausal encounter, the same method of adding relative pose and splitting the queues used in the causal join is used to generate the causal and acausal virtual robots created by this encounter.

The following figure explains the joining process graphically. The scenario starts with three robots x^1 , x^2 , and x^3 , with x^1 being the sole robot in the initial mapping set. The actual robots are represented as white circles, the causal particles as red circles, the acausal particles as blue circles, and encounters as purple arrows. At timestep T_1 , an encounter occurs between x^2 and x^3 , but neither is in the mapping set at this time, so this encounter is stored. At timestep 2, robots x^1 and x^2 encounter each other, spawning a causal and acausal pair of pose particles x^{2+} and x^{2-} respectively. x^{1+} and x^{2+} continue propagating forward with new data (not shown) while x^{2-} propagates backwards in time using stored odometry and measurements. The stored encounter between x^2 and x^3 that was previously stored is then incorporated at time T_3 when the acausal particles x^{2-} reach the location of the encounter. This spawns a set of particles for the previously encounter robot x^3 . As the original encounter happened on the first step, no acausal particles are spawned at this time. As x^{2-} reaches the last remaining pose in the acausal queue, the particles are removed as the past data is exhausted.

3) *FastSLAM Updates*: At each time-step, the pose posterior is updated for all pose states within the current mapping set are updated according to a modified Rao-Blackwellized particle filter. This process first applies an odometry prediction to each particle. This single update applies to each causal robot particle a forward odometry change given an arbitrary non-linear and noisy forward odometry model $g(x, u)$. For each virtual robot traversing its path in reverse, the next stored point is popped off the corresponding stack and used with a noisy inverse odometry model $g_r(x, \bar{u})$

$$\begin{aligned} x_{t+1}^{r(i)} &= g(x_t^{r(i)}, u_{t+1}^{r(i)}) \\ \bar{x}_{t+1}^{r(i)} &= g_r(\bar{x}_t^{r(i)}, \bar{u}_{t+1}^{r(i)}) \end{aligned}$$

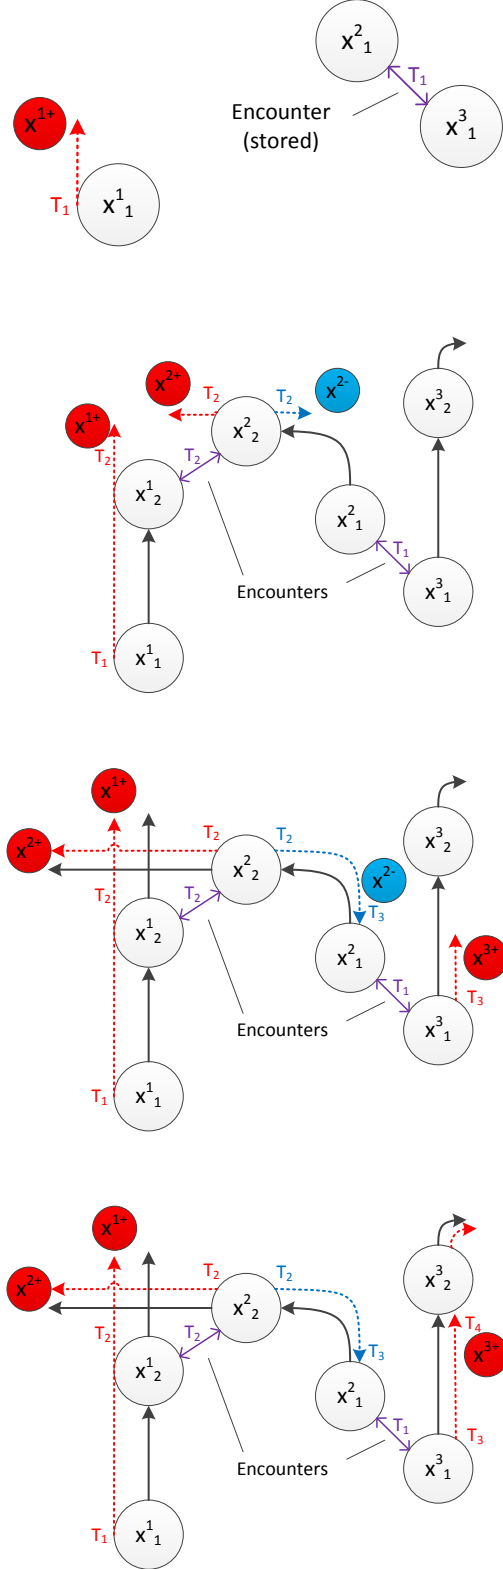




Fig. 2. Example of Robot Encounters in this MRSLAM Implementation

IV. EXPERIMENTAL VALIDATION

$$\begin{bmatrix} x_t \\ y_t \theta_t \end{bmatrix} = \begin{bmatrix} x_{t-1} \\ y_{t-1} \theta_{t-1} \end{bmatrix} + \begin{bmatrix} dx_t \cos(\theta_t + d\theta_t) \\ dy_t \sin(\theta_t + d\theta_t) \\ \theta_{t-1} + d\theta_t \end{bmatrix} \quad (1)$$



(a) The test geometry, robot paths, and encounters (gray).



(b) Robot encounters at a given time.

Fig. 3.

V. CONCLUSION

In this paper we have presented Howard's algorithm from [?] for the integration of data from several independent robots, and demonstrated a successful implementation. Further, we have modified the algorithm to use a particle filter for each robot and demonstrated that it can qualitatively

REFERENCES

- [1] Tim Bailey and Hugh Durrant-Whyte. Simultaneous localization and mapping (slam): Part ii. *IEEE Robotics & Automation Magazine*, 13(3):108–117, 2006.
- [2] Andreas Birk and Stefano Carpin. Merging occupancy grid maps from multiple robots. *Proceedings of the IEEE*, 94(7):1384–1397, 2006.
- [3] Wolfram Burgard, Mark Moors, Cyrill Stachniss, and Frank E Schneider. Coordinated multi-robot exploration. *Robotics, IEEE Transactions on*, 21(3):376–386, 2005.
- [4] Hugh Durrant-Whyte and Tim Bailey. Simultaneous localization and mapping: part i. *Robotics & Automation Magazine, IEEE*, 13(2):99–110, 2006.
- [5] Dieter Fox, Jonathan Ko, Kurt Konolige, Benson Limketkai, Dirk Schulz, and Benjamin Stewart. Distributed multirobot exploration and mapping. *Proceedings of the IEEE*, 94(7):1325–1339, 2006.
- [6] Andrew Howard. Multi-robot simultaneous localization and mapping using particle filters. *The International Journal of Robotics Research*, 25(12):1243–1256, 2006.
- [7] Andrew Howard and Nicholas Roy. The robotics data set repository (radish), 2003.
- [8] Sebastian Thrun. A probabilistic on-line mapping algorithm for teams of mobile robots. *The International Journal of Robotics Research*, 20(5):335–363, 2001.